



PERCONA
Performance Consulting Experts

How to choose High Availability solutions for MySQL

MySQL UC 2010

Yves Trudeau

Read by Peter Zaitsev

Percona Inc

MySQLPerformanceBlog.com

About us

- <http://www.percona.com>
- <http://www.mysqlperformanceblog.com/>
- <http://www.bigdbahead.com>

Yves Trudeau, Ph. D.
Principal Consultant

Plan

- 1) Definitions of some High-Availability (HA) terms
- 2) Questions to ask
- 3) HA mindset
- 4) Common HA solutions with MySQL
 - Replication based
 - Shared storage based
 - NDB Cluster
- 5) Other solutions

Definitions

1) High-Availability (HA)

- A computer architecture design and implementation that is targeted at improving the availability of a given service

2) Uptime and downtime

- The proportion of time a high availability service is up or down over the total time. Normally, uptime + downtime = 100%.

3) Level of availability

- Typically in term of the fraction of uptime and referred by the number of 9, 99% (2 9s), 99.9% (3 9s), etc.

Definitions

4) Single point of failure (SPOF)

- An isolated device or piece of software for which a failure will cause a downtime of the HA service. The goal of an HA architecture is to remove the SPOFs.

5) Recovering or failover

- The process by which a HA architecture recovers after a failure.

6) Fencing/Stonith

- Often, an HA architecture is stuck by a non-responsive device that is not releasing a critical resource. Fencing or Stonith (*Shoot The Other Node In The Head*) is then required.

Definitions

7) Cluster

- A group of computers acting together to offer a service.

8) Fault Tolerance

- Ability to handle failures with graceful degradation. Not all components may need same level of HA

9) Disaster Recovery

- The plan and technologies to recover in case of disaster. Often longer downtime allowed in this case.

Questions

1) Do you need HA?

- can be rephrased to “What is your downtime cost?”
- Include non-monetary aspects like corporate image and marketing
- For the downtime cost, what is acceptable over a year?
- Do you have maintenance windows that offers reduce downtime cost?

2) Can you afford to lose some data?

- What is the cost of losing a transaction?
- How critical is data consistency ?

Questions

3) Are relying on MyISAM only features?

- Fulltext indexes?
- GIS?
- Sphinx or Lucene options?

4) What is the write load?

- How many threads are writing simultaneously?
- How many write ops/s?

5) What is the growth potential of your dataset?

Questions

- 6) How qualified is your IT department or support company?
- 7) How much are you ready to invest?

HA Mindset

1) HA, not only about technologies

- No technology is fool proof
- Operating procedures are required
- Testing and staging
- Monitoring and alerting

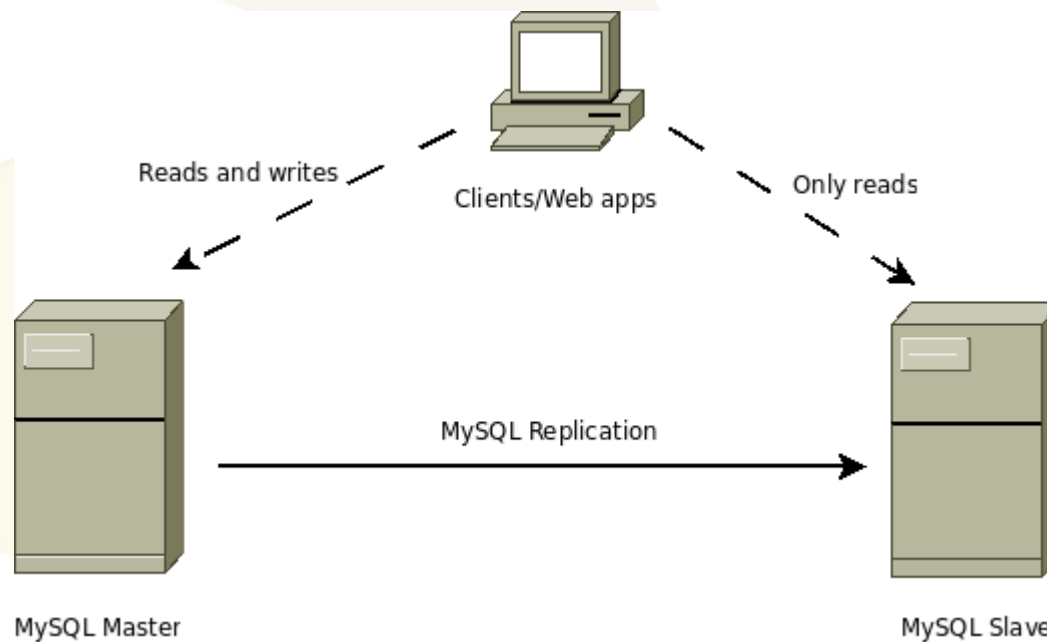
2) A HA is not isolated, look at the broad picture

- No need for HA of 99.999% if ISP SLA is 99.9%
- Power
- Cooling, more frequent problem than you might think
- Very high HA requirements need multiple data centers.

Replication based

1) Simplest example, plain replication

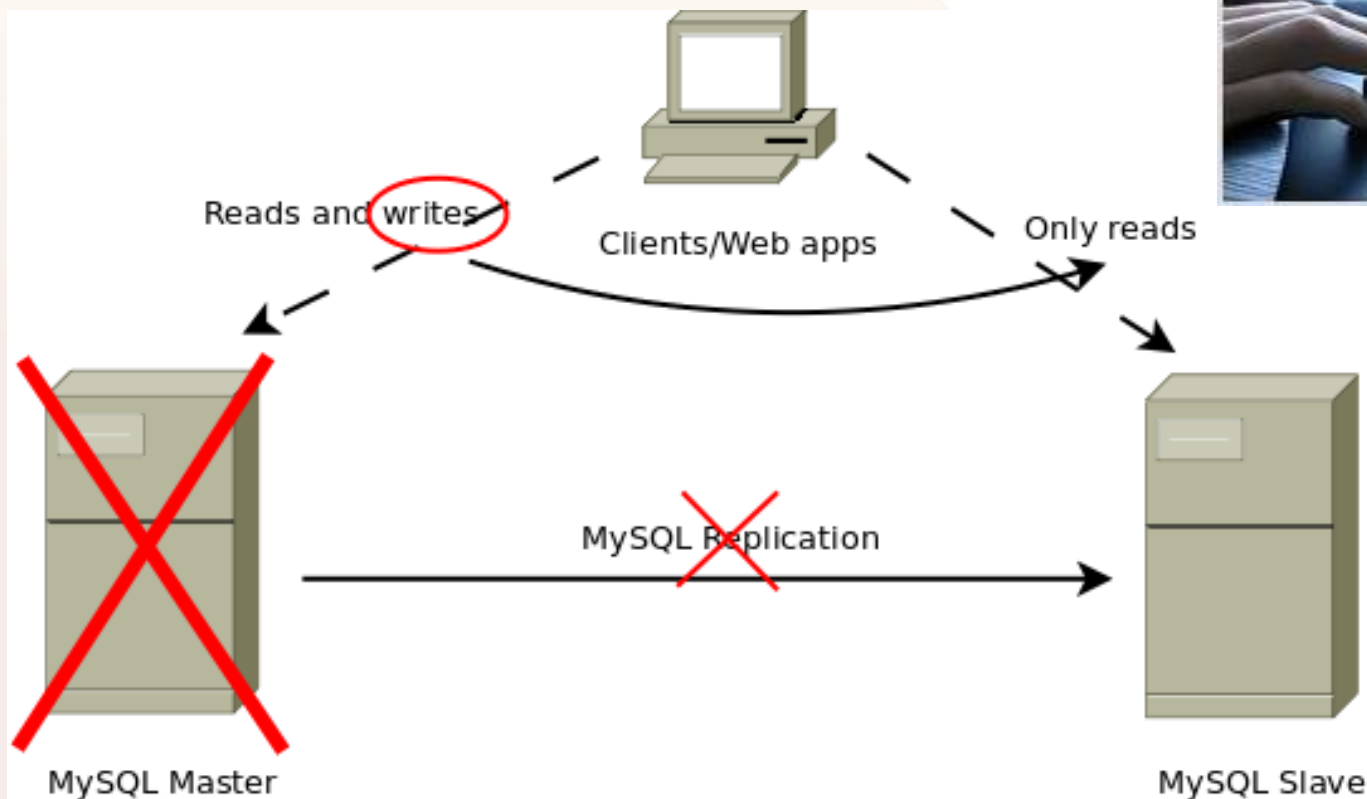
- Widely used
- Manual failover



Replication based failover

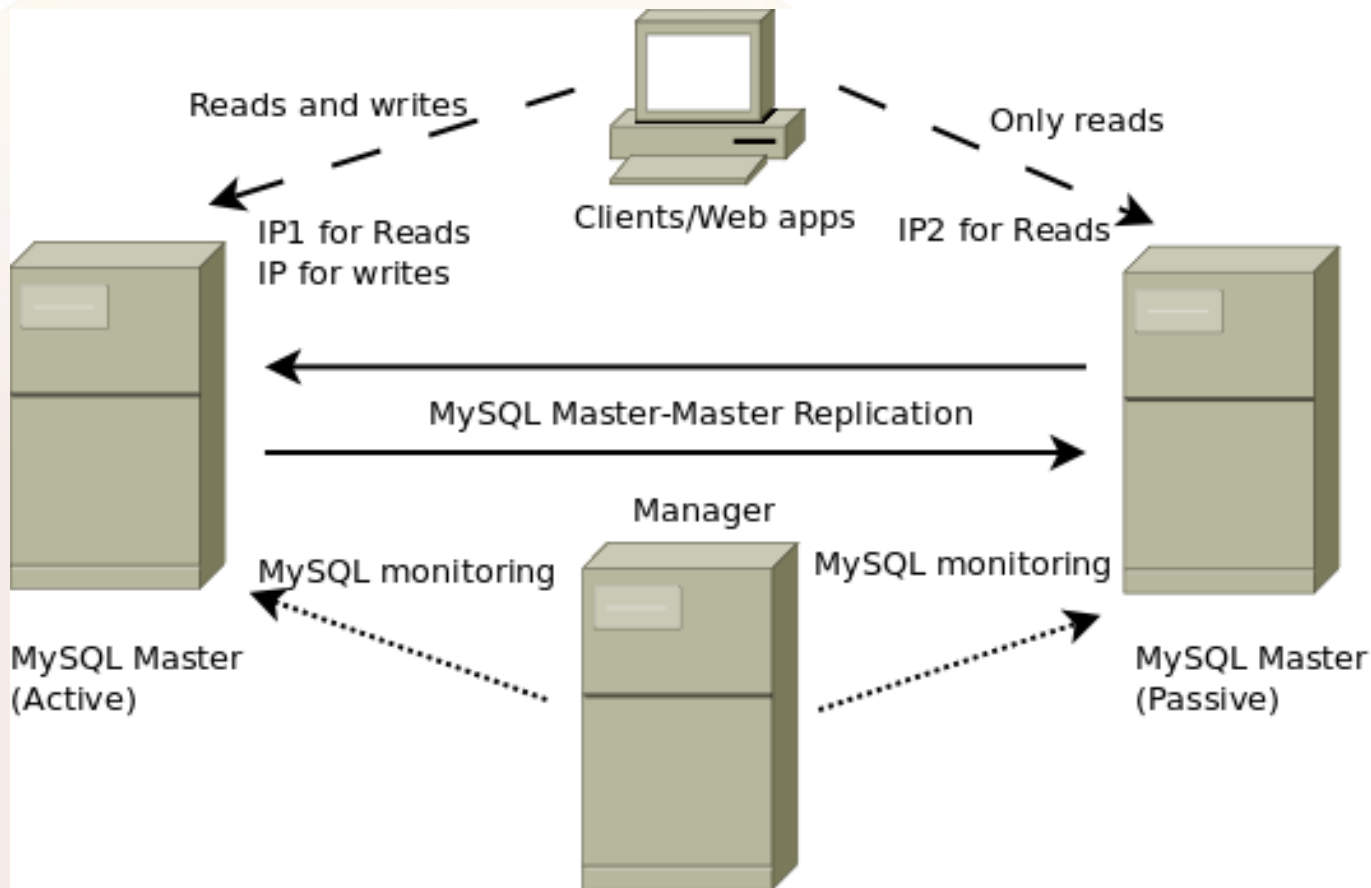
2) Simple replication, failover process

- Manual operation required



Replication based MMM

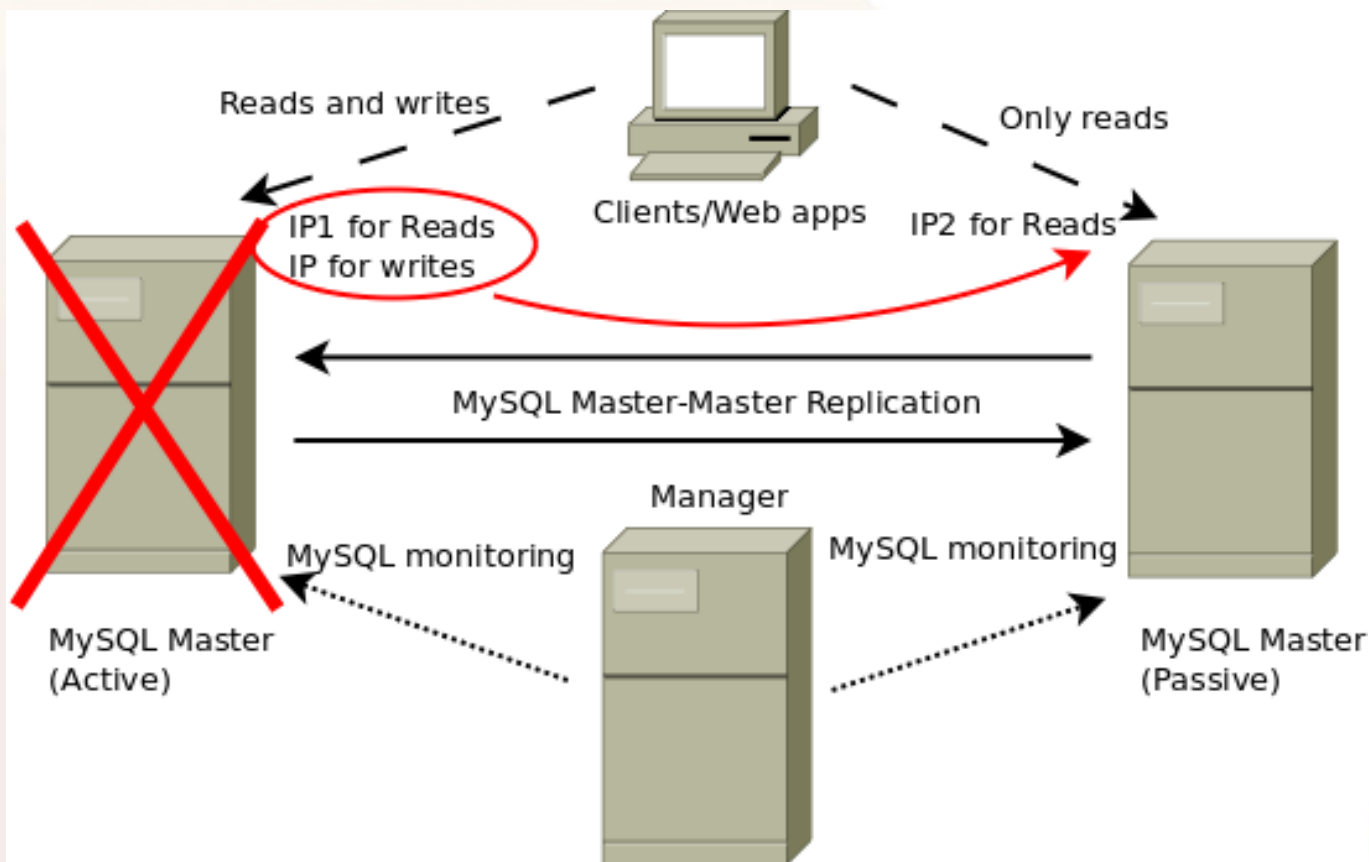
3) Example 2, using MMM



Replication based MMM failover

4) Failover with MMM

- Manager transfer IP1 and IP to the surviving server



Replication based other

4) Other solutions built on replication

- Tungsten, Java proxy layer doing man in the middle work for queries and replication stream
- Pacemaker/Heartbeat, not released yet, developed by Linbit, will add fencing capabilities


Replication based Pros

- Simple/Inexpensive
- Supports MyISAM
- All servers can be used, no standby
- Good to scale read ops
- Caches are kept warm
- Can be used for online schema changes, upgrades
- Loosely coupled

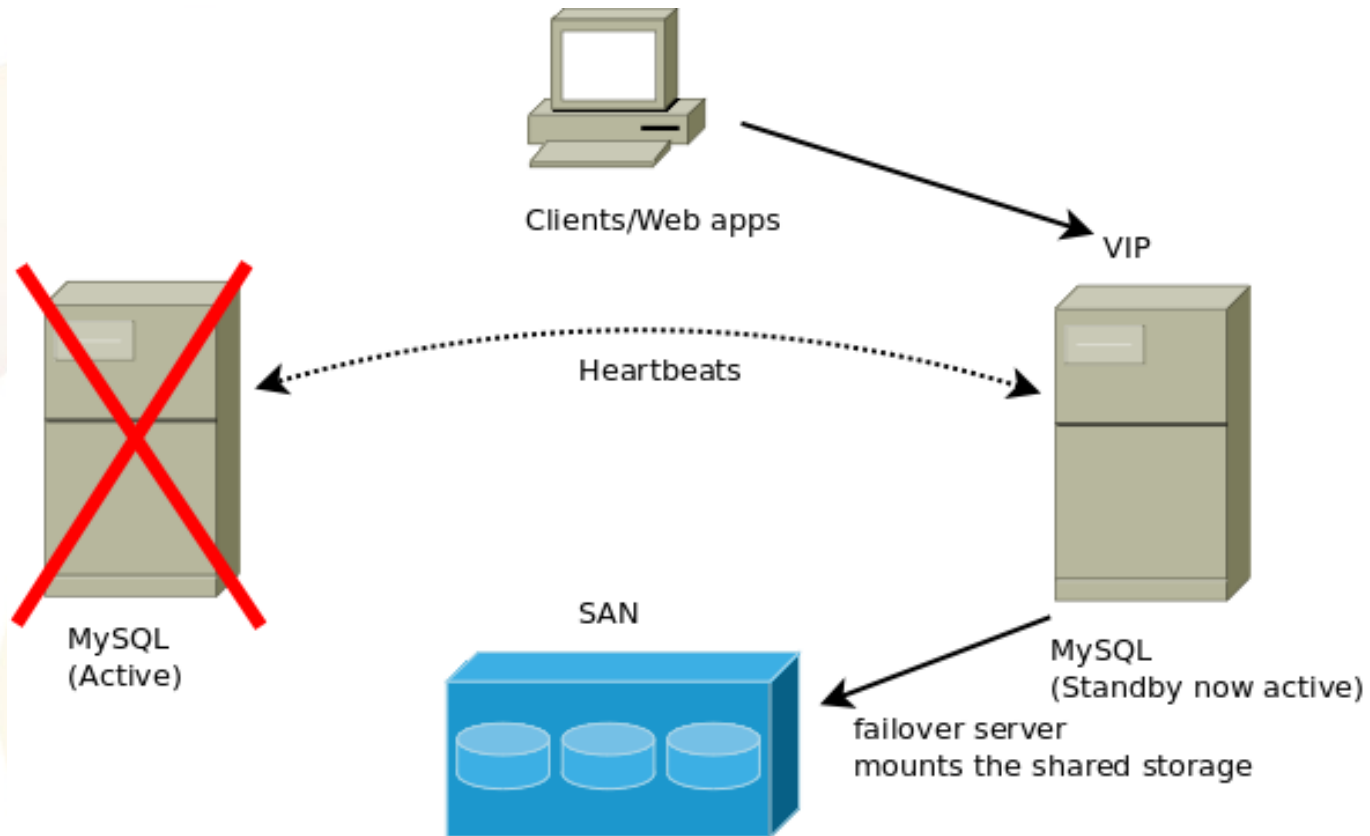
Replication based Cons

- Limited availability
 - Replication can break
 - Replication can lag behind
 - Replication can be out of sync
- Manual or at best semi-automatic failover, tricky to automate.
- Limited write capacity: single threaded
- Can lose data: async (with semi-sync repl?)
- Immature tools, edge cases not always handled

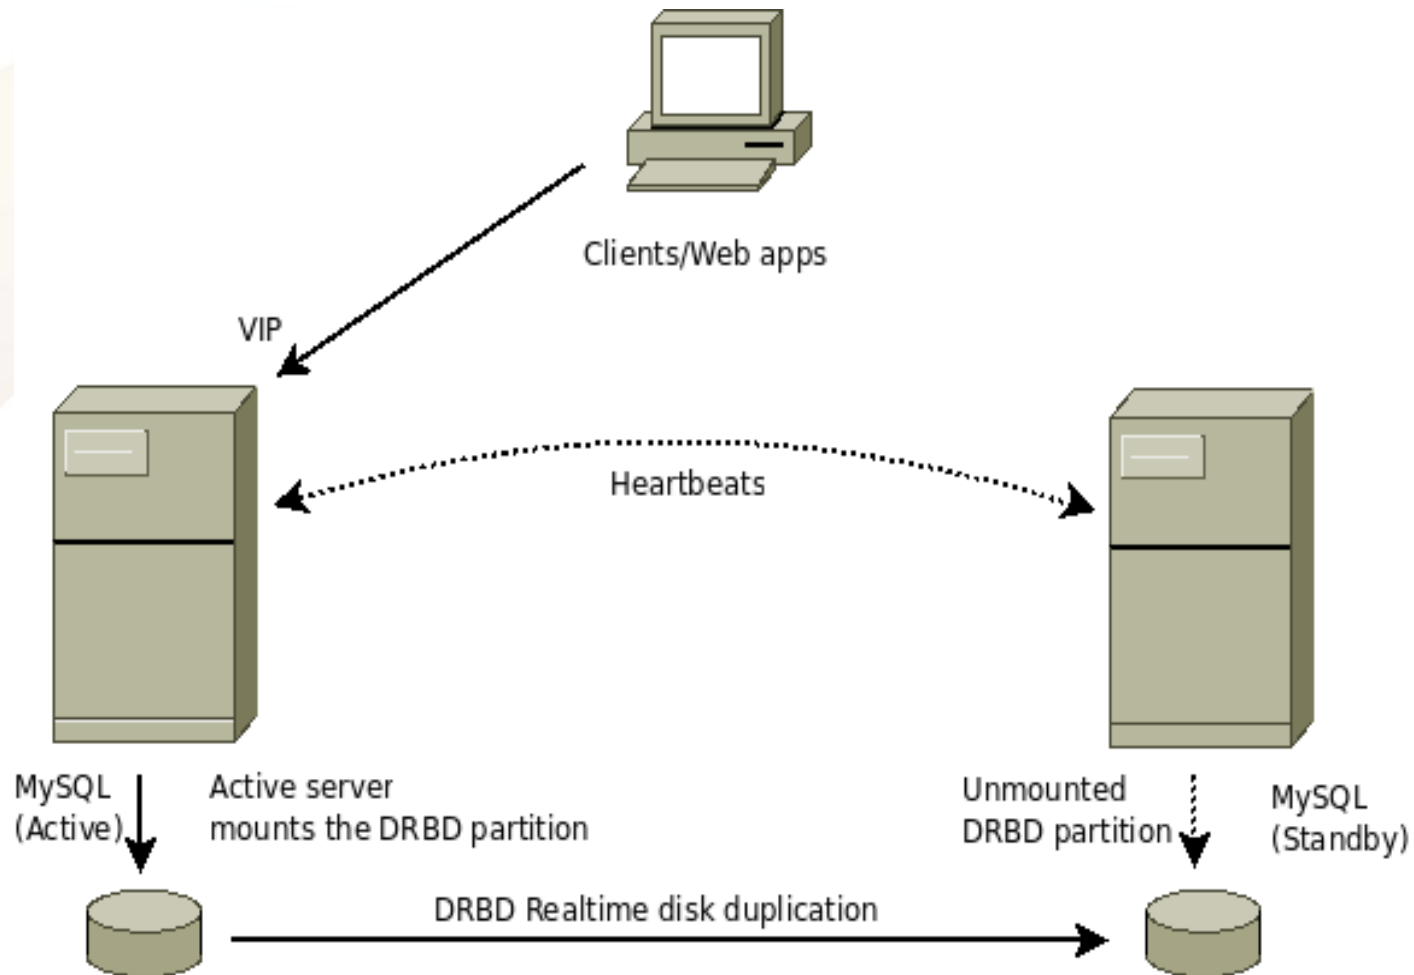
Shared storage/SAN



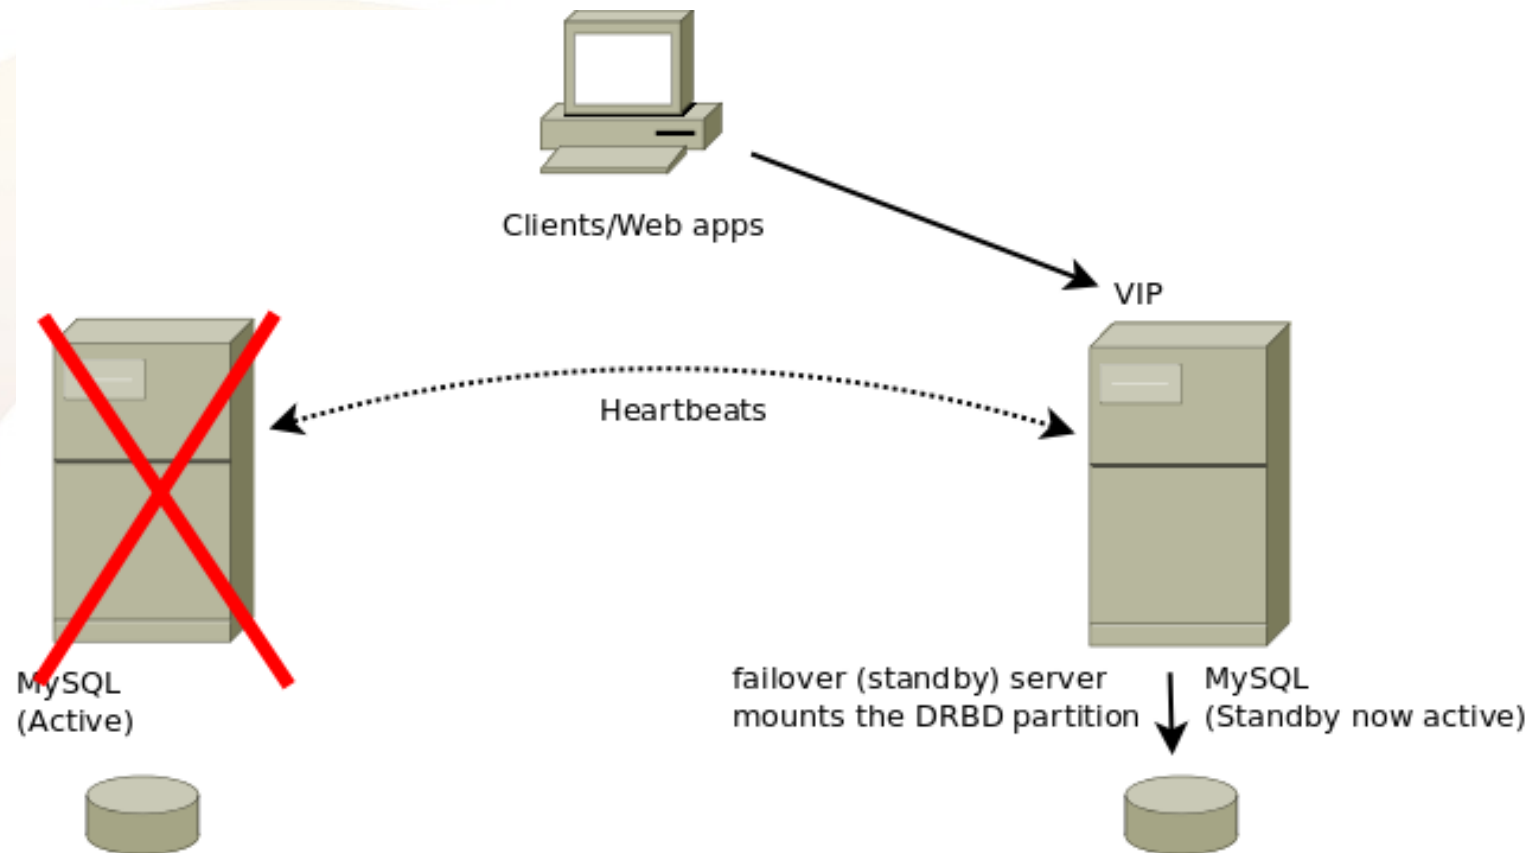
Shared storage/SAN failover



Shared storage/DRBD



Shared storage/DRBD failover



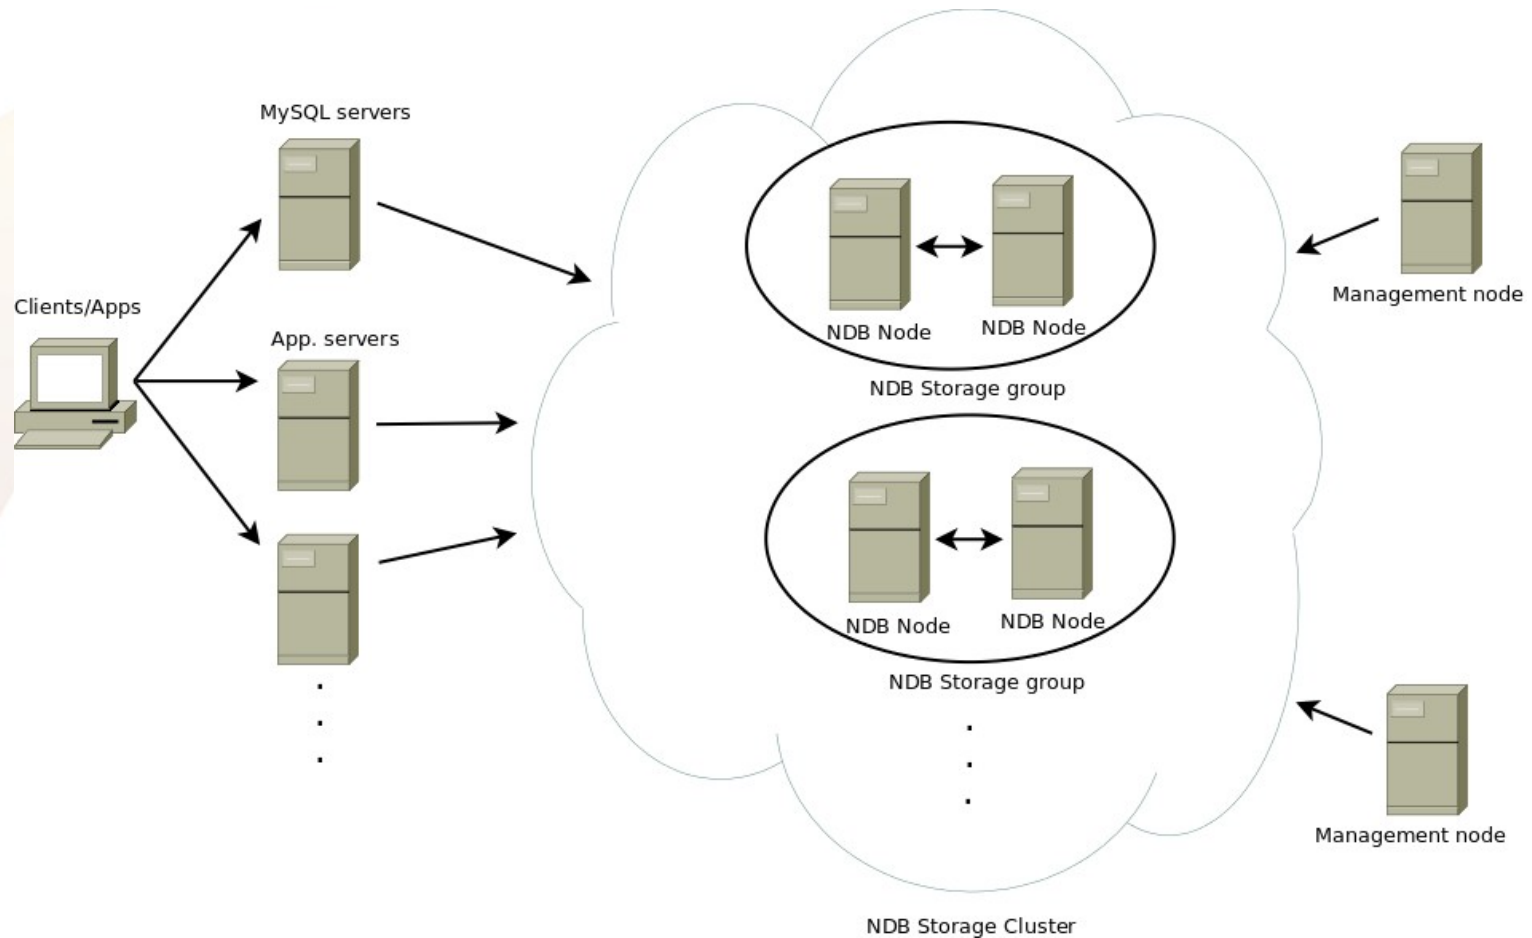
Shared storage Pros

- No data loss
- Much higher write capacity
- Automatic failover in about 1 minute with InnoDB log files of about 100 MB
 - Comes at performance cost
- No SPOF with DRBD

Shared storage Cons

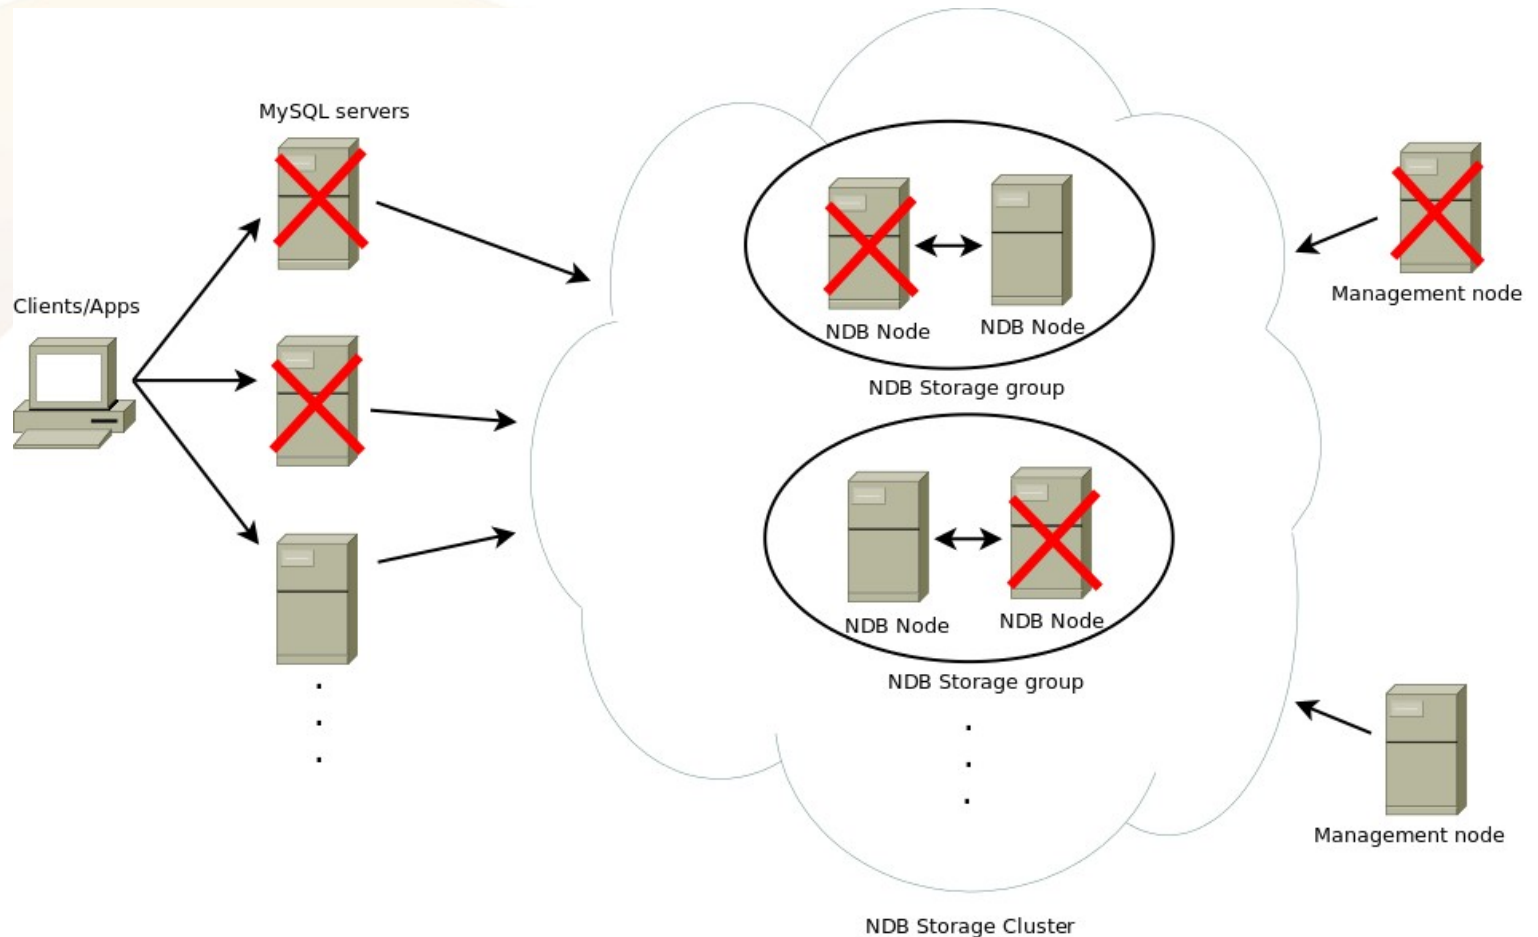
- Only works with engine supporting recovery (InnoDB), should work with PBXT and Maria (Have not tested)
- More complex: nic bounding, fencing, etc.
- Requires fencing
- A server is standby, idle hardware
- Cold cache after failover although XtraDB LRU dump can be a big winner here
- No online schema change
- Corruption Propagation
-

NDB Cluster



NDB Cluster failover

Still up!



NDB Cluster Pros

- Sharding framework by hashes
- No SPOF, high level of HA
- Scalable, if the schema is well designed, adding data nodes adds processing capacity
- Huge write capacity

NDB Cluster Cons

- Complex, much than other solutions
- Needs work on schema and queries for good performance
- Higher skill set required
- Poor for large joins
- Size of dataset more limited, large memory footprint
- Minimum of physical servers

Emerging Solutions

- 1) Hot market a lot of work going on !
- 2) ScaleDB
- 3) Galera Replication
- 4) Number of Solutions which are still in stealth

Mixed solutions

- Geo-redundancy
 - Shared storage + replication
 - NDB + replication
- Scaling reads ops
 - Shared storage HA Master with reads slaves
- Sharding solutions
 - Global database using NDB
 - shards using shared storage

Questions

Questions?

Yves Trudeau: yves@percona.com

<http://www.percona.com/>

<http://www.mysqlperformanceblog.com/>